

PROBLEM STATEMENT:- TO PREDICT THE RAIN FALL BASED ON VARIOUS FEATURES OF THE DATASET

```
In [1]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\Y.Saranya\Downloads\rainfall in india 1901-2015.csv")
df
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NO
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.

4116 rows × 19 columns



DATA PREPROCESSING

In [3]: `df.head()`

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	D
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	16
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	22
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	4
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	34

In [4]: `df.tail()`

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0

```
In [5]: df.isnull().any()
```

```
Out[5]: SUBDIVISION    False
        YEAR          False
        JAN           True
        FEB           True
        MAR           True
        APR           True
        MAY           True
        JUN           True
        JUL           True
        AUG           True
        SEP           True
        OCT           True
        NOV           True
        DEC           True
        ANNUAL        True
        Jan-Feb       True
        Mar-May       True
        Jun-Sep       True
        Oct-Dec       True
        dtype: bool
```

```
In [6]: df.fillna(method='ffill',inplace=True)
        df.isnull().sum()
```

```
Out[6]: SUBDIVISION    0
        YEAR          0
        JAN           0
        FEB           0
        MAR           0
        APR           0
        MAY           0
        JUN           0
        JUL           0
        AUG           0
        SEP           0
        OCT           0
        NOV           0
        DEC           0
        ANNUAL        0
        Jan-Feb       0
        Mar-May       0
        Jun-Sep       0
        Oct-Dec       0
        dtype: int64
```

In [7]: `df.describe()`

Out[7]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4
mean	1958.218659	18.957240	21.823251	27.415379	43.160641	85.788994	230.567979	
std	33.140898	33.576192	35.922602	47.045473	67.816588	123.220150	234.896056	
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000	
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.475000	
50%	1958.000000	6.000000	6.700000	7.900000	15.700000	36.700000	138.900000	
75%	1987.000000	22.200000	26.800000	31.400000	50.125000	97.400000	306.150000	
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000	2

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN            4116 non-null   float64
3   FEB            4116 non-null   float64
4   MAR            4116 non-null   float64
5   APR            4116 non-null   float64
6   MAY            4116 non-null   float64
7   JUN            4116 non-null   float64
8   JUL            4116 non-null   float64
9   AUG            4116 non-null   float64
10  SEP            4116 non-null   float64
11  OCT            4116 non-null   float64
12  NOV            4116 non-null   float64
13  DEC            4116 non-null   float64
14  ANNUAL         4116 non-null   float64
15  Jan-Feb       4116 non-null   float64
16  Mar-May       4116 non-null   float64
17  Jun-Sep       4116 non-null   float64
18  Oct-Dec       4116 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [9]: `df.columns`

Out[9]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May', 'Jun-Sep', 'Oct-Dec'], dtype='object')

```
In [10]: df.shape
```

```
Out[10]: (4116, 19)
```

```
In [11]: df['ANNUAL'].value_counts()
```

```
Out[11]: 790.5      4
         770.3      4
         1836.2     4
         1024.6     4
         1926.5     3
         ..
         443.9      1
         689.0      1
         605.2      1
         509.7      1
         1642.9     1
         Name: ANNUAL, Length: 3712, dtype: int64
```

```
In [12]: df['Jan-Feb'].value_counts()
```

```
Out[12]: 0.0      238
         0.1       80
         0.2       52
         0.3       38
         0.4       32
         ...
         23.3        1
         95.2        1
         76.9        1
         66.5        1
         69.3        1
         Name: Jan-Feb, Length: 1220, dtype: int64
```

```
In [13]: df['Mar-May'].value_counts()
```

```
Out[13]: 0.0      29
         0.1      13
         0.3      11
         8.3      11
         11.5     10
         ..
         246.3     1
         248.1     1
         151.3     1
         249.5     1
         223.9     1
         Name: Mar-May, Length: 2262, dtype: int64
```

```
In [14]: df['Jun-Sep'].value_counts()
```


```
Out[14]: 434.3      4
          334.8      4
          573.8      4
          613.3      4
          1082.3     3
          ..
          301.6      1
          380.9      1
          409.3      1
          229.4      1
          958.5      1
          Name: Jun-Sep, Length: 3683, dtype: int64
```

```
In [15]: df['Oct-Dec'].value_counts()
```

```
Out[15]: 0.0      16
          0.1      15
          0.5      13
          0.6      12
          0.7      11
          ..
          191.5     1
          124.5     1
          139.1     1
          41.5      1
          555.4     1
          Name: Oct-Dec, Length: 2389, dtype: int64
```

EXPLORATORY DATA ANALYSIS

```
In [16]: df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
In [17]: df.columns
```

```
Out[17]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```

```
In [18]: x=df[["FEB"]]
y=df[["JAN"]]
```

LINEAR REGRESSION

```
In [19]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [20]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

9.650666612303553

Out[20]:


	coefficient
FEB	0.442278

```
In [21]: score=reg.score(X_test,y_test)
print(score)
```


0.1793580786264921

```
In [22]: predictions=reg.predict(X_test)
plt.scatter(y_test,predictions)
```

Out[22]: <matplotlib.collections.PathCollection at 0x23bfa68f070>




```
In [23]: df500=df[:][:500]
sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)
plt.show()
```




```
In [24]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
reg.fit(X_train,y_train)
reg.fit(X_test,y_test)
```

```
Out[24]: LinearRegression
LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [28]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='black')
plt.plot(X_test,y_pred,color='red')
plt.show()
```



```
In [29]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.22901901017995652

Ridge model

```
In [33]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [34]: features= df.columns[0:5]
target= df.columns[-5]
x=np.array(df['JAN']).reshape(-1,1)
y=np.array(df['FEB']).reshape(-1,2)
x= df[features].values
y= df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
```

```
In [35]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```

```
In [36]: print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

the train score for ridge model is0.9999999999874192
the test score for ridge model is0.99999999998833

```
In [37]: lr=LinearRegression()
```

LASSO MODEL


```
In [39]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.9999207747038827
The test score for ls model is0.9999206791315255

```
In [40]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

```
Out[40]: <Axes: >
```



```
In [41]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.9999999999999921
0.9999999999999921
```

ELASTIC NET

```
In [43]: from sklearn.linear_model import ElasticNet
eln=ElasticNet()
eln.fit(x,y)
print(eln.coef_)
print(eln.intercept_)
print(eln.score(x,y))
```

```
[9.99098574e-01 0.00000000e+00 3.02728910e-05 0.00000000e+00
 0.00000000e+00]
0.016258606966616185
0.9999992160905338
```

In []: